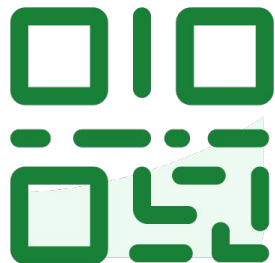




slido



Join at [slido.com](https://slido.com)  
#2777102

① Click **Present with Slido** or install our [Chrome extension](#) to display joining instructions for participants while presenting.



## LECTURE 2

# Pandas, Part I

Introduction to **pandas** syntax, operators, and functions

**Data 100/Data 200, Spring 2025 @ UC Berkeley**

Narges Norouzi and Josh Grossman



# Goals for This Lecture

---

Lecture 02, Data 100 Spring 2025

- Introduce **pandas**, an important Python library for working with data
- Key data structures: DataFrames, Series, Indices
- Extracting data: `loc`, `iloc`, `[]`

This is the first of a three-lecture sequence about **pandas**.

Get ready: lots of code incoming!

- Lecture: introduce high-level concepts
- Lab, homework: practical experimentation



# Agenda

---

Lecture 02, Data 100 Spring 2025

- Tabular data
- Series, DataFrames, and Indices
- Data extraction with `loc`, `iloc`, and `[]`



# Tabular Data


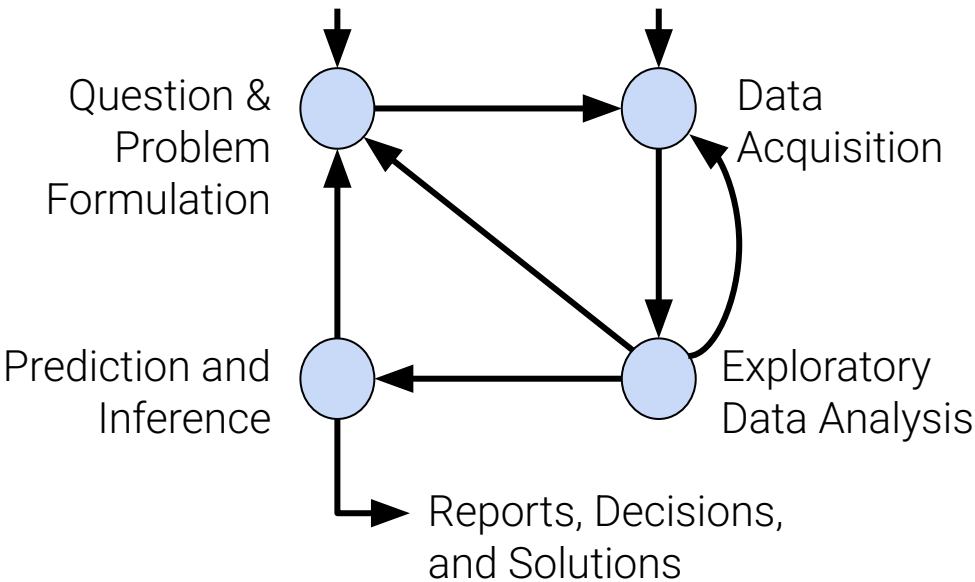
---


Lecture 02, Data 100 Spring 2025

- **Tabular data**
- Series, DataFrames, and Indices
- Data extraction with `loc`, `iloc`, and `[]`



?





(Weeks 1 and 2)

Exploring and Cleaning Tabular Data  
From **datascience** to **pandas**

(Weeks 2 and 3)

Data Science in Practice  
EDA, Data Cleaning, Text processing (regular expressions)



## Congratulations!!!

You **have collected** or **have been given** a box of data.

What does this "data" actually look like?  
How will you work with it?





"Tabular data" = data in a table.

Typically:

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

A **row** represents one observation (here, a single person running for president in a particular year).

A **column** represents some characteristic, or feature, of that observation (here, the political party of that person).

In Data 8, you worked with the datascience library using Tables.  
In Data 100 (and beyond), we'll use an industry-standard library called pandas.





The Python Data  
Analysis Library



Stands for "panel  
data"

The (unofficial) Data  
100 logo



a cartoon panda



Using pandas, we can:

- Arrange data in a tabular format.
- Extract useful information filtered by specific conditions.
- Operate on data to gain new insights.
- Apply NumPy functions to our data (our friends from Data 8).
- Perform vectorized computations to speed up our analysis (Lab 1).

pandas is the standard tool across research and industry for working with tabular data.

The first two weeks of Data 100 will serve as a "bootcamp" in helping you build familiarity with operating on data with pandas.

Your Data 8 knowledge will serve you well! Much of our work will be in [translating syntax](#).



📁 / ... / lecture / lec02 /

Name ▲

📁 data

Data used in this lecture

📖 data8\_translation\_e...

Unofficial **datascience** -> **pandas** translations

📖 lec02.ipynb

**Primary notebook for lecture**



# DataFrames, Series, and Indices

---

Lecture 02, Data 100 Spring 2025

- Tabular data
- **DataFrames, Series, and Indices**
- Data extraction with `loc`, `iloc`, and `[]`



# DataFrames

In the "language" of pandas, we call a table a **DataFrame**.

We think of **DataFrames** as collections of named columns, called **Series**.

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

0      Andrew Jackson

1      John Quincy Adams

2      Andrew Jackson

3      John Quincy Adams

4      Andrew Jackson

...

182     Donald Trump

183     Kamala Harris

184     Jill Stein

185     Robert Kennedy

186     Chase Oliver

Name: Candidate, Length: 187, dtype: object

A DataFrame

A Series named "Candidate"



# Series

A **Series** is a 1-dimensional array-like object. It contains:

- A sequence of **values** of the same type.
- A sequence of data labels, called the **index**.

**pd** is the conventional alias for **pandas**

```
import pandas as pd  
s = pd.Series(["welcome", "to", "data 100"])
```

0	welcome
1	to
2	data 100

dtype: object

**Index**, accessed by calling `s.index`

```
RangeIndex(start=0, stop=3, step=1)
```

**Values**, accessed by calling `s.values`

```
array(['welcome', 'to', 'data 100'], dtype=object)
```



- We can provide index labels for items in a **Series** by passing an index list.

```
s = pd.Series([-1, 10, 2], index = ["a", "b", "c"])
```

```
a    -1  
b    10  
c     2  
dtype: int64
```

```
s.index
```

```
Index(['a', 'b', 'c'], dtype='object')
```

- A **Series** index can also be changed.

```
s.index = ["first", "second", "third"]
```

```
first    -1  
second   10  
third     2  
dtype: int64
```

```
s.index
```

```
Index(['first', 'second', 'third'], dtype='object')
```





- We can select a single value or a set of values in a **Series** using:
  - A single label
  - A list of labels
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a      4
b     -2
c      0
d      6
dtype: int64
```



- We can select a single value or a set of values in a **Series** using:
  - **A single label**
  - A list of labels
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a    4
b   -2
c    0
d    6
dtype: int64
```

```
s["a"]
```

```
4
```



## Selection in Series

- We can select a single value or a set of values in a **Series** using:
  - A single label
  - **A list of labels**
  - A filtering condition

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a    4
b   -2
c    0
d    6
dtype: int64
```

```
s[["a", "c"]]
```

```
a    4
c    0
dtype: int64
```



## Selection in Series

- We can select a single value or a set of values in a **Series** using:
  - A single label
  - A list of labels
  - **A filtering condition**

```
s = pd.Series([4, -2, 0, 6], index = ["a", "b", "c", "d"])
```

```
a      4
b     -2
c      0
d      6
dtype: int64
```

- Say we want to select values in the **Series** that satisfy a particular condition:
  - 1) Apply a boolean condition to the **Series**. This creates a **new Series of boolean values**.
  - 2) Index into our **Series** using this boolean condition. **pandas** will select only the entries in the **Series** that satisfy the condition.

```
s > 0
```

```
a      True
b     False
c     False
d      True
dtype: bool
```

```
s[s > 0]
```

```
a      4
d      6
dtype: int64
```



slido



# What is the output of the following code?

① Click **Present with Slido** or install our [Chrome extension](#) to activate this poll while presenting.



2777102

# DataFrames of Series!

Typically, we will work with **Series** using the perspective that they are columns in a **DataFrame**.

We can think of a **DataFrame** as a collection of **Series** that all share the same **Index**.

0	1824	0	Andrew Jackson
1	1824	1	John Quincy Adams
2	1828	2	Andrew Jackson
3	1828	3	John Quincy Adams
4	1832	4	Andrew Jackson



...	...	...	...
182	2024	182	Donald Trump
183	2024	183	Kamala Harris
184	2024	184	Jill Stein
185	2024	185	Robert Kennedy
186	2024	186	Chase Oliver

Name: Year,

Name: Candidate,

[...]



	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

The Series "Year"

The Series "Candidate"

The DataFrame **elections**



The syntax of creating **DataFrame** is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- From a dictionary.
- From a **Series**.



The syntax of creating `DataFrame` is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a `DataFrame`. Here, we will go over the most popular ones.

- **From a CSV file.**
- Using a list and column name(s).
- From a dictionary.
- From a **Series**.

```
elections = pd.read_csv("data/elections.csv")
```

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

The DataFrame `elections`





The syntax of creating `DataFrame` is:

```
pandas.DataFrame(data, index, columns)
```

Many approaches exist for creating a `DataFrame`. Here, we will go over the most popular ones.

- **From a CSV file.** `elections = pd.read_csv("data/elections.csv", index_col="Year")`
- Using a list and column name(s).
- From a dictionary.
- From a `Series`.

	Candidate	Party	Popular vote	Result	%
Year					
1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
1828	Andrew Jackson	Democratic	642806	win	56.203927
1828	John Quincy Adams	National Republican	500897	loss	43.796073
1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...
2024	Donald Trump	Republican	77303568	win	49.808629
2024	Kamala Harris	Democratic	75019230	loss	48.336772
2024	Jill Stein	Green	861155	loss	0.554864
2024	Robert Kennedy	Independent	756383	loss	0.487357
2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

The `DataFrame` `elections` with `"Year"` as `Index`



2777102

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- **Using a list and column name(s).**
- From a dictionary.
- From a **Series**.

```
pd.DataFrame([1, 2, 3],  
             columns=["Numbers"])
```

Numbers	
0	1
1	2
2	3

```
pd.DataFrame([[1, "one"], [2, "two"]],  
             columns = ["Number", "Description"])
```

Number Description		
0	1	one
1	2	two



2777102

## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- **From a dictionary.**
- From a **Series**.

Specify columns of the **DataFrame**

```
pd.DataFrame({"Fruit":["Strawberry", "Orange"],  
             "Price": [5.49, 3.99]})
```

```
pd.DataFrame([{"Fruit":"Strawberry", "Price":5.49},  
             {"Fruit":"Orange", "Price":3.99}])
```

Specify rows of the **DataFrame**

	Fruit	Price
0	Strawberry	5.49
1	Orange	3.99



## Creating a DataFrame

Many approaches exist for creating a **DataFrame**. Here, we will go over the most popular ones.

- From a CSV file.
- Using a list and column name(s).
- From a dictionary.
- **From a Series.**

```
s_a = pd.Series(["a1", "a2", "a3"], index = ["r1", "r2", "r3"])  
s_b = pd.Series(["b1", "b2", "b3"], index = ["r1", "r2", "r3"])
```

```
pd.DataFrame({"A-column":s_a, "B-column":s_b})
```

	A-column	B-column
r1	a1	b1
r2	a2	b2
r3	a3	b3

```
pd.DataFrame(s_a)
```

```
s_a.to_frame()
```

	0
r1	a1
r2	a2
r3	a3



# Indices Are Not Necessarily Row Numbers

An **Index** (a.k.a. row labels) can also:

- Be non-numeric.
- Have a name, e.g. "Candidate".

```
# Creating a DataFrame from a CSV file and specifying the Index column  
elections = pd.read_csv("data/elections.csv", index_col = "Candidate")
```

Candidate	Year	Party	Popular vote	Result	%
Andrew Jackson	1824	Democratic-Republican	151271	loss	57.210122
John Quincy Adams	1824	Democratic-Republican	113142	win	42.789878
Andrew Jackson	1828	Democratic	642806	win	56.203927
John Quincy Adams	1828	National Republican	500897	loss	43.796073
Andrew Jackson	1832	Democratic	702735	win	54.574789



2777102

## Indices Are Not Necessarily Unique

The row labels that constitute an index do not have to be unique.

- Left: The **index** values are all unique and numeric, acting as a row number.
- Right: The **index** values are named and non-unique.

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win

	Candidate	Party	%	Result
Year				
2008	Obama	Democratic	52.9	win
2008	McCain	Republican	45.7	loss
2012	Obama	Democratic	51.1	win
2012	Romney	Republican	47.2	loss
2016	Clinton	Democratic	48.2	loss
2016	Trump	Republican	46.1	win



- We can select a new column and set it as the index of the **DataFrame**.

Example: Setting the index to the "Candidate" column.

```
elections.set_index("Candidate")
```

	Year	Party	Popular vote	Result	%
Candidate					
Andrew Jackson	1824	Democratic-Republican	151271	loss	57.210122
John Quincy Adams	1824	Democratic-Republican	113142	win	42.789878
Andrew Jackson	1828	Democratic	642806	win	56.203927
John Quincy Adams	1828	National Republican	500897	loss	43.796073
Andrew Jackson	1832	Democratic	702735	win	54.574789
...	...	...	...	...	...
Donald Trump	2024	Republican	77303568	win	49.808629
Kamala Harris	2024	Democratic	75019230	loss	48.336772
Jill Stein	2024	Green	861155	loss	0.554864
Robert Kennedy	2024	Independent	756383	loss	0.487357
Chase Oliver	2024	Libertarian Party	650130	loss	0.418895



2777102

## Resetting the Index

- We can change our mind and reset the **Index** back to the default list of integers.

`elections.reset_index()`

	Year	Party	Popular vote	Result	%
Candidate					
Andrew Jackson	1824	Democratic-Republican	151271	loss	57.210122
John Quincy Adams	1824	Democratic-Republican	113142	win	42.789878
Andrew Jackson	1828	Democratic	642806	win	56.203927
John Quincy Adams	1828	National Republican	500897	loss	43.796073
Andrew Jackson	1832	Democratic	702735	win	54.574789
...	...	...	...	...	...
Donald Trump	2024	Republican	77303568	win	49.808629
Kamala Harris	2024	Democratic	75019230	loss	48.336772
Jill Stein	2024	Green	861155	loss	0.554864
Robert Kennedy	2024	Independent	756383	loss	0.487357
Chase Oliver	2024	Libertarian Party	650130	loss	0.418895

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895







2777102

## Column Names Are Usually Unique!

Column names in **pandas** are almost always unique.

- Example: Really shouldn't have two columns named "Candidate".

	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win



Sometimes you'll want to extract the list of row and column labels.

```
elections.set_index("Party")
```

For row labels, use `DataFrame.index`:

```
elections.index
```

```
Index(['Democratic-Republican', 'Democratic-Republican', 'Democratic',  
      'National Republican', 'Democratic', 'National Republican',  
      'Anti-Masonic', 'Whig', 'Democratic', 'Whig',  
      ...  
      'Green', 'Democratic', 'Republican', 'Libertarian', 'Green',  
      'Republican', 'Democratic', 'Green', 'Independent',  
      'Libertarian Party'],  
      dtype='object', name='Party', length=187)
```

For column labels, use `DataFrame.columns`:

```
elections.columns
```

```
Index(['Candidate', 'Year', 'Popular vote', 'Result', '%'], dtype='object')
```

For shape of the `DataFrame` we use `DataFrame.shape`:

```
elections.shape
```

```
(187, 6)
```




# The Relationship Between DataFrames, Series, and Indices

We can think of a **DataFrame** as a collection of **Series** that all share the same **Index**.

- Candidate, Party, %, Year, and Result **Series** all share an **Index** from 0 to 5.

Candidate Series    Party Series    % Series    Year Series    Result Series



	Candidate	Party	%	Year	Result
0	Obama	Democratic	52.9	2008	win
1	McCain	Republican	45.7	2008	loss
2	Obama	Democratic	51.1	2012	win
3	Romney	Republican	47.2	2012	loss
4	Clinton	Democratic	48.2	2016	loss
5	Trump	Republican	46.1	2016	win



slido



# Which of the following lines of code creates this DataFrame?

① Click **Present with Slido** or install our [Chrome extension](#) to activate this poll while presenting.



The API for the **DataFrame** class is enormous.

- API: "Application Programming Interface".
- The API is the set of abstractions supported by the class.

Full documentation is at

<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.html>

- Compare with the **Table** class from Data8: <http://data8.org/datascience/tables.html>
- We will only consider a tiny portion of this API.

We want you to get familiar with the real world programming practice of... Googling!

- Answers to your questions are often found in the **pandas** documentation, Stack Overflow, etc.



2777102

# Interlude

---

Break!





# Data Extraction with `loc`, `iloc`, and `[]`

---

Lecture 02, Data 100 Spring 2025

- Tabular data
- DataFrames, Series, and Indices
- **Data extraction with `loc`, `iloc`, and `[]`**



One of the most basic tasks for manipulating a **DataFrame** is to extract rows and columns of interest. As we'll see, the large **pandas** API means there are many ways to do things.

Common ways we may want to extract data:

- Grab the first or last **n** rows in the **DataFrame**.
- Grab data with a certain label.
- Grab data at a certain position.

We'll find that all three of these methods are useful to us in data manipulation tasks.





2777102

## .head and .tail

The simplest scenarios: We want to extract the first or last **n** rows from the **DataFrame**.

- `df.head(n)` will return the first **n** rows of the DataFrame **df**.
- `df.tail(n)` will return the last **n** rows.

### elections

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

### elections.head(5)

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789

### elections.tail(5)

	Year	Candidate	Party	Popular vote	Result	%
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895



## Label-based Extraction: `.loc`

A more complex task: We want to extract data with specific column or index labels.

```
df.loc[row_labels, column_labels]
```

The `.loc` accessor allows us to specify the **labels** of rows and columns we wish to extract.

- We describe "labels" as the bolded text at the top and left of a **DataFrame**.

	Year	Candidate	Party	Popular vote	Result	%
<b>0</b>	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
<b>1</b>	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
<b>2</b>	1828	Andrew Jackson	Democratic	642806	win	56.203927
<b>3</b>	1828	John Quincy Adams	National Republican	500897	loss	43.796073
<b>4</b>	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
<b>182</b>	2024	Donald Trump	Republican	77303568	win	49.808629
<b>183</b>	2024	Kamala Harris	Democratic	75019230	loss	48.336772
<b>184</b>	2024	Jill Stein	Green	861155	loss	0.554864
<b>185</b>	2024	Robert Kennedy	Independent	756383	loss	0.487357
<b>186</b>	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895

Row labels

Column labels



Arguments to `.loc` can be:

- A list.
- A slice (syntax is inclusive of the right hand side of the slice).
- A single value.

	Year	Candidate	Party	Popular vote	Result	%
<b>0</b>	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
<b>1</b>	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
<b>2</b>	1828	Andrew Jackson	Democratic	642806	win	56.203927
<b>3</b>	1828	John Quincy Adams	National Republican	500897	loss	43.796073
<b>4</b>	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
<b>182</b>	2024	Donald Trump	Republican	77303568	win	49.808629
<b>183</b>	2024	Kamala Harris	Democratic	75019230	loss	48.336772
<b>184</b>	2024	Jill Stein	Green	861155	loss	0.554864
<b>185</b>	2024	Robert Kennedy	Independent	756383	loss	0.487357
<b>186</b>	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895



## Label-based Extraction: .loc

Arguments to `.loc` can be:

- **A list.**
- A slice (syntax is inclusive of the right hand side of the slice).
- A single value.

```
elections.loc[[87, 25, 179], ["Year", "Candidate", "Result"]]
```

Select the rows with labels 87, 25, and 179.

	Year	Candidate	Result
87	1932	Herbert Hoover	loss
25	1860	John C. Breckinridge	loss
179	2020	Donald Trump	loss

Select the columns with labels "Year", "Candidate", and "Result".



## Label-based Extraction: .loc

Arguments to `.loc` can be:

- A list.
- **A slice** (syntax is **inclusive of the right hand side of the slice**).
- A single value.

```
elections.loc[[87, 25, 179], "Popular vote": "%"]
```

Select the rows with  
labels 87, 25, and 179.

	Popular vote	Result	%
87	15761254	loss	39.830594
25	848019	loss	18.138998
179	74216154	loss	46.858542

Select all columns *starting*  
from "Popular vote" *until* "%".



## Label-based Extraction: .loc

To extract *all* rows or *all* columns, use a colon (:)

```
elections.loc[:, ["Year", "Candidate", "Result"]]
```

All rows for the columns with labels "Year", "Candidate", and "Result".

Ellipses (...) indicate more rows not shown. →

```
elections.loc[[87, 25, 179], :]
```

All columns for the rows with labels 87, 25, 179.

	Candidate	Year	Party	Popular vote	Result	%
87	Herbert Hoover	1932	Republican	15761254	loss	39.830594
25	John C. Breckinridge	1860	Southern Democratic	848019	loss	18.138998
179	Donald Trump	2020	Republican	74216154	loss	46.858542

	Year	Candidate	Result
0	1824	Andrew Jackson	loss
1	1824	John Quincy Adams	win
2	1828	Andrew Jackson	win
3	1828	John Quincy Adams	loss
4	1832	Andrew Jackson	win
...	...	...	...
182	2024	Donald Trump	win
183	2024	Kamala Harris	loss
184	2024	Jill Stein	loss
185	2024	Robert Kennedy	loss
186	2024	Chase Oliver	loss



## Label-based Extraction: .loc

Arguments to `.loc` can be:

- A list.
- A slice (syntax is inclusive of the right hand side of the slice).
- **A single value.**

```
elections.loc[[87, 25, 179], "Popular vote"]
```

```
87      15761254
```

```
25       848019
```

```
179     74216154
```

```
Name: Popular vote, dtype: int64
```

Wait, what? Why did everything get so ugly?

We've extracted a subset of the "Popular vote" column as a **Series**.

```
elections.loc[0, "Candidate"]
```

```
'Andrew Jackson'
```

We've extracted the string value with row label 0 and column label "Candidate".



# Lecture 2 ended here!

We will cover the rest in lecture 3





## Integer-based Extraction: `.iloc`

A different scenario: We want to extract data according to its *position*.

- Example: Grab the 1st, 2nd, and 3rd columns of the **DataFrame**.

```
df.iloc[row_integers, column_integers]
```

The `.iloc` accessor allows us to specify the **integers** of rows and columns we wish to extract.

- Python convention: The first position has integer index 0.

		0	1	2	3	4	5	Column integers
		Year	Candidate	Party	Popular vote	Result	%	
Row integers	0	0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
	1	1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
	2	2	1828	Andrew Jackson	Democratic	642806	win	56.203927
	3	3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
	4	4	1832	Andrew Jackson	Democratic	702735	win	54.574789
		...	...	...	...	...	...	...



Arguments to `.iloc` can be:

- A list.
- A slice (syntax is **exclusive** of the right hand side of the slice).
- A single value.

	Year	Candidate	Party	Popular vote	Result	%
0	1824	Andrew Jackson	Democratic-Republican	151271	loss	57.210122
1	1824	John Quincy Adams	Democratic-Republican	113142	win	42.789878
2	1828	Andrew Jackson	Democratic	642806	win	56.203927
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
...	...	...	...	...	...	...
182	2024	Donald Trump	Republican	77303568	win	49.808629
183	2024	Kamala Harris	Democratic	75019230	loss	48.336772
184	2024	Jill Stein	Green	861155	loss	0.554864
185	2024	Robert Kennedy	Independent	756383	loss	0.487357
186	2024	Chase Oliver	Libertarian Party	650130	loss	0.418895



## Integer-based Extraction: `.iloc`

Arguments to `.iloc` can be:

- **A list.**
- A slice (syntax is **exclusive** of the right hand side of the slice).
- A single value.

```
elections.iloc[[1, 2, 3], [0, 1, 2]]
```

Select the rows at  
positions 1, 2, and 3.

	Year	Candidate	Party
1	1824	John Quincy Adams	Democratic-Republican
2	1828	Andrew Jackson	Democratic
3	1828	John Quincy Adams	National Republican

Select the columns  
at positions 0, 1,  
and 2.



## Integer-based Extraction: `.iloc`

Arguments to `.iloc` can be:

- A list.
- **A slice** (syntax is **exclusive of the right hand side of the slice**).
- A single value.

```
elections.iloc[[1, 2, 3], 0:3]
```

Select the rows  
at positions 1, 2,  
and 3.

	Year	Candidate	Party
1	1824	John Quincy Adams	Democratic-Republican
2	1828	Andrew Jackson	Democratic
3	1828	John Quincy Adams	National Republican



Select *all* columns from  
integer 0 to integer 2.

Remember:  
integer-based slicing is  
right-end exclusive!



## Integer-based Extraction: `.iloc`

Just like `.loc`, we can use a colon with `.iloc` to extract all rows or all columns.

```
elections.iloc[:, 0:3]
```

	Year	Candidate	Result
0	1824	Andrew Jackson	loss
1	1824	John Quincy Adams	win
2	1828	Andrew Jackson	win
3	1828	John Quincy Adams	loss
4	1832	Andrew Jackson	win
...	...	...	...
182	2024	Donald Trump	win
183	2024	Kamala Harris	loss
184	2024	Jill Stein	loss
185	2024	Robert Kennedy	loss
186	2024	Chase Oliver	loss

Grab all rows of the columns at integers 0 to 2.



## Integer-based Extraction: `.iloc`

Arguments to `.iloc` can be:

- A list.
- A slice (syntax is exclusive of the right hand side of the slice).
- **A single value.**

```
elections.iloc[[1, 2, 3], 1]
```

```
1    John Quincy Adams
2      Andrew Jackson
3    John Quincy Adams
Name: Candidate, dtype: object
```

As before, the result for a single value argument is a **Series**.

We have extracted row integers 1, 2, and 3 from the column at position 1.

```
elections.iloc[0, 1]
```

```
'Andrew Jackson'
```

We've extracted the string value with row position 0 and column position 1.



## `.loc` vs `.iloc`

Remember:

- `.loc` performs **label-based** extraction
- `.iloc` performs **integer-based** extraction

When choosing between `.loc` and `.iloc`, you'll usually choose `.loc`.

- Safer: If the order of data gets shuffled in a public database, your code still works.
- Readable: Easier to understand what `elections.loc[:, ["Year", "Candidate", "Result"]]` means than `elections.iloc[:, [0, 1, 4]]`

`.iloc` can still be useful.

- Example: If you have a **DataFrame** of movie earnings sorted by earnings, can use `.iloc` to get the median earnings for a given year (index into the middle).







Selection operators:

- **.loc** selects items by **label**. First argument is rows, second argument is columns.
- **.iloc** selects items by **integer**. First argument is rows, second argument is columns.
- **[ ]** only takes one argument, which may be:
  - A slice of **row numbers**.
  - A list of **column labels**.
  - A single **column label**.

That is, **[ ]** is context sensitive.

Let's see some examples.



2777102

## Context-dependent Extraction: [ ]

[ ] only takes one argument, which may be:

- **A slice of row integers.**
- A list of column labels.
- A single column label.

elections[3:7]

	Year	Candidate	Party	Popular vote	Result	%
3	1828	John Quincy Adams	National Republican	500897	loss	43.796073
4	1832	Andrew Jackson	Democratic	702735	win	54.574789
5	1832	Henry Clay	National Republican	484205	loss	37.603628
6	1832	William Wirt	Anti-Masonic	100715	loss	7.821583



2777102

## Context-dependent Extraction: [ ]

[ ] only takes one argument, which may be:

- A slice of row numbers.
- **A list of column labels.**
- A single column label.

```
elections[["Year", "Candidate", "Result"]]
```

	Year	Candidate	Result
0	1824	Andrew Jackson	loss
1	1824	John Quincy Adams	win
2	1828	Andrew Jackson	win
3	1828	John Quincy Adams	loss
4	1832	Andrew Jackson	win
...	...	...	...
182	2024	Donald Trump	win
183	2024	Kamala Harris	loss
184	2024	Jill Stein	loss
185	2024	Robert Kennedy	loss
186	2024	Chase Oliver	loss



2777102

## Context-dependent Extraction: [ ]

[ ] only takes one argument, which may be:

- A slice of row numbers.
- A list of column labels.
- **A single column label.**

```
elections["Candidate"]
```

```
0      Andrew Jackson
1      John Quincy Adams
2      Andrew Jackson
3      John Quincy Adams
4      Andrew Jackson
```

```
...
```

```
182     Donald Trump
183     Kamala Harris
184       Jill Stein
185    Robert Kennedy
186     Chase Oliver
```

```
Name: Candidate, Length: 187, dtype: object
```

Extract the "Candidate" column as a **Series**.



In short: [] can be much more concise than `.loc` or `.iloc`

- Consider the case where we wish to extract the "Candidate" column. It is far simpler to write `elections["Candidate"]` than it is to write `elections.loc[:, "Candidate"]`

In practice, [] is often used over `.iloc` and `.loc` in data science work. Typing time adds up!

slido



Which of the following statements correctly returns the value "blue fish" from the "weird" DataFrame?

- ① Click **Present with Slido** or install our [Chrome extension](#) to activate this poll while presenting.



## LECTURE 2

# Pandas, Part I

Content credit: [Acknowledgments](#)